

Prompt Definition — Member Playbook

Stage: Prompt Definition (the first stage) **Goal of this doc:** help you get your prompt signed off on the first try. **Version:** 1.0 · **Last updated:** 30 Jun 2026

1. What this stage is, in one minute

You take a real task from your own work and write it as a single instruction for an AI agent (Codex). The agent should be able to read that instruction, go off, use a couple of real tools, make some judgment calls, and come back with something actually done — not just a document.

A reviewer then reads your prompt and either signs it off or sends it back with comments. This playbook is about clearing the reviewer's bar the first time, so you skip the back-and-forth.

The whole trick: **be exact about *what* you want and *where* things live, and stay open about *how* the agent gets there.**

2. The bar your prompt has to clear

These are the five things a reviewer is checking for. Hit all five and you're most of the way home.

1. Real work across at least two tools. The agent has to read or pull information from one place, think about it, and then *do something* in another place — update a sheet, draft an email, create a calendar event, open a ticket, post a message. If the only output is a report, doc, or summary, it doesn't pass. Making a document can be *part* of the task, just not the whole thing.

2. One focused task. Something you'd genuinely hand to an agent in real life and let it run — like "review this week's leads and draft the follow-ups," not "run my whole sales quarter." If it spans weeks or your entire role, it's too big.

3. Hard enough to be interesting. When you test it, Codex should score around **1-4 out of 7**. If it scores 5-7, the task was too easy and the reviewer will ask you to add difficulty. Aim for something with real judgment, trade-offs, or messy inputs.

4. Nothing confidential. No employer or customer secrets, no private personal data, no passwords or keys. If your real workflow touches any of that, swap in a personal or made-up stand-in from the start (e.g., a personal Google Sheet instead of a company system).

5. Exact about workspace and outputs, open about method. Name the actual folder, file, sheet, tab, channel, or repo. Spell out what each output should contain. But **don't** write step 1, step 2, step 3 — let the agent figure out the approach. You set the destination and the rules; the agent drives.

3. The Workflow Prompt Maker

This is a prompt you can use to shape your real workflow into a clean, review-ready one. Fill in the four blanks at the top, then let it produce the prompt. It's written to pull the structure *out of your own task* — so what comes out is your real workflow, just sharpened.

Tip: the more honest detail you put in the four blanks, the better the result.
Garbage in, garbage out.

Workflow Prompt Maker — copy this and fill in the four blanks

You help me turn a real task from my job into one workflow prompt for an AI agent
↪ (Codex).

Write the final prompt in plain, human language – short sentences, no buzzwords, and no numbered step-by-step instructions. I tell the agent WHAT I want and WHERE things are; the agent decides HOW.

MY WORKFLOW (what I actually do, start to finish):
[describe it like you'd explain it to a new teammate]

TOOLS I CAN GIVE THE AGENT (only apps I can personally log into and authorize):
[e.g., Gmail, a specific Google Sheet, Google Calendar, GitHub, Notion]

WHERE THE INFO LIVES (be exact – real names):
[folder / file / sheet / tab / label / channel names]

WHERE THE AGENT SHOULD ACT (the exact destination it writes to or changes):
[which sheet+tab, which draft, which calendar, which ticket or channel]

Now write ONE workflow prompt that:

- Reads from at least one named source and takes a real action in at least one other named tool – not just producing a document.
- Names the exact files, folders, sheets, tabs, channels, or repos. No "the tracker" or "a spreadsheet."
- Says exactly what each output should contain – the fields, columns, or format.
- Uses exact dates and names the timezone. No "today," "this week," or "recently."
- Gives any limits, thresholds, or budgets as real numbers, and says what to do when two options are basically tied.
- Says what must already exist or be reachable before it starts, and to stop and tell me if something's missing or it can't get access.
- For anything that sends, posts, or shares outside, says whether to draft it for my review or send it, and to whom.
- Spells out what "done well" looks like – the success criteria.
- Says what to do if data is missing or a step fails.
- Stays to one focused task, not my whole month of work.
- Leaves out anything confidential – use personal or made-up stand-ins instead.
- States the goal, the sources, and the rules, then lets the agent choose its own approach. Do not write step 1 / step 2 / step 3.

Keep it tight and readable, the way I'd explain it to a sharp new teammate.

What “filling the blanks” looks like

My workflow: Every Monday I go through last week's customer feedback, find the ones that need a follow-up, log them, and send the customer a reply.

Tools: Gmail, Google Sheet “Customer Issue Tracker”, a Google Drive folder of feedback.

Where info lives: Drive folder “Q2 Customer Feedback”; Gmail label “support”.

Where the agent acts: the “Customer Issue Tracker” sheet, and Gmail drafts.

What the maker gives you back (roughly):

Using the feedback docs in the Google Drive folder “Q2 Customer Feedback” and the emails under the Gmail label “support” from 16–22 Jun 2026 (IST), find the recurring, high-priority issues. Update the Google Sheet “Customer Issue Tracker” so each issue has a category, priority, owner, customer impact, and a recommended next action. Where two issues look like the same root problem, merge them rather than logging duplicates. Then prepare Gmail **drafts** (don’t send) to customers marked “High Priority,” each including the issue, the proposed fix, and the next step we’re asking them to take. If you can’t open the folder, the label, or the sheet, stop and tell me which one. Treat the work as done when every high-priority issue is in the tracker with all five fields and has a matching draft.

Notice what it does: names real places, exact dates with timezone, says what each output contains, handles duplicates, drafts instead of sends, checks access first, and defines “done.” That’s a first-pass prompt.

4. Pre-submit checklist

Run down this list before you submit. Each item is something reviewers actually flag. If you can tick all of them, you’re in good shape. (Ordered by how often it trips people up.)

- ☐ **Exact dates + timezone.** No “today” or “this week.” Write real dates and say IST/PST/etc.
- ☐ **No vague words.** Anywhere you wrote “relevant,” “appropriate,” “the file,” or “as needed,” replace it with the specific thing.
- ☐ **Named sources.** Every place the agent reads from has a real name and, if it matters, a time window.
- ☐ **Output spelled out.** You listed the exact fields/columns/format each output must have.
- ☐ **Named destinations.** The exact sheet+tab / draft / calendar / channel / repo the agent writes to — not a generic app name.
- ☐ **Send vs. draft is clear.** For anything leaving your workspace (email, message, post), you said whether to draft for review or send, and to whom.
- ☐ **Starting point + access check.** You said what must already exist and told the agent to stop and report if it can’t reach something.
- ☐ **Privacy handled.** No confidential or personal data; sensitive actions are called out.
- ☐ **Checkpoints / rubric.** If the task has a “looks right” bar, you named it.
- ☐ **Two-plus tools, real action.** It reads somewhere and acts somewhere — not document-only.
- ☐ **Numbers are exact.** Limits, budgets, thresholds are real numbers, with a tie-breaker rule.
- ☐ **Only tools you can grant.** Every app named is one you can actually log into and authorize.
- ☐ **Success is defined.** There’s a clear “this is done when...”
- ☐ **Missing-data fallback.** You said what to do if something’s absent or fails.
- ☐ **Format named.** If output format matters, you said what it should be.

- ❑ **Metadata is realistic.** Your time-per-run and frequency look believable.
- ❑ **No duplicates / safe to re-run.** If run twice, it won't create double entries.

5. The most common reasons prompts get sent back — and the fix

These are the top failure modes pulled from real reviewer feedback across the program. Each one is “what the reviewer says” → “what to change.”

Fuzzy dates and times. *“Replace ‘recent’ with exact dates and a timezone.”* → Write 16–22 Jun 2026 (IST), not “last week.”

Vague wording. *“This reference is ambiguous.”* → Name the exact file, person, or rule instead of “the doc” or “the relevant items.”

Unclear sources. *“Define where the data comes from.”* → List the exact source(s) and the slice of data (which folder, which label, which dates).

No input/output schema. *“What should the output contain?”* → Spell out the fields or columns and the format.

Unnamed tools or destinations. *“Which sheet? Which channel?”* → Give the real names: sheet + tab, #channel, org/repo.

No confirmation step for external actions. *“Should this send or just draft?”* → Say “draft for my review” or “send to X,” explicitly.

Missing starting state. *“What needs to exist first?”* → State the prerequisites and tell the agent to stop and report if access is missing.

Privacy gaps. *“Remove restricted data / require authorization.”* → Strip confidential inputs; flag any sensitive action.

No checkpoints. *“How does it know it’s on track?”* → Add a simple rubric or checkpoint.

Not really multi-tool. *“This is mostly producing a document.”* → Add a real action in a second app.

Soft thresholds. *“Give exact numbers and a tie-breaker.”* → “under \$500 per person; if two options tie, pick the cheaper.”

Tools you can’t grant. *“You can’t authorize this app.”* → Swap in a personal equivalent you can actually connect.

No success criteria. *“What does success look like?”* → Add a one-line “done when...”

No error fallback. *“What if the data isn’t there?”* → Say what to do when something’s missing.

Over-engineered method. *“Don’t script every step.”* → Delete the step-by-step; state the goal and constraints and let the agent choose.

6. A worked example: before → after

Before (would get bounced):

Use the customer feedback in Google Drive and the tracker to send follow-up emails to priority customers.

Why it fails: which Drive folder? which tracker? what fields? send or draft? what dates? It’s agentic in spirit but vague on everything that matters.

After (first-pass ready):

Using the feedback docs in Google Drive folder “Q2 Customer Feedback” and the Google Sheet “Customer Issue Tracker,” find recurring high-priority issues from 16–22 Jun 2026 (IST). Update the tracker so each issue has category, priority, owner, customer impact, and recommended next action; merge duplicates rather than re-logging them. Then create Gmail **drafts** (don’t send) to customers marked “High Priority,” each with the issue, the proposed fix, and the next step. If you can’t open the folder, the sheet, or Gmail, stop and tell me which. Done when every high-priority issue is in the tracker with all five fields and has a matching draft.

Same task. The difference is named places, exact dates, defined outputs, a draft-not-send rule, an access check, and a clear finish line.

7. Metadata that matters

When you fill in the task fields, a few of them quietly affect your review:

- **Difficulty / Codex rating:** you’re aiming for a task Codex finds genuinely hard (scores 1–4 of 7). If it breezes through, plan to add a wrinkle — more sources, a trade-off, messier data.
- **Apps / connectors:** list the exact brand names (e.g., “Gmail, Google Sheets, GitHub”), not “email” or “a spreadsheet.”
- **Time per run and frequency:** keep them realistic. A 5-minute task you do 200 times a month, or a 12-hour one-off, both look off.
- **Occupation / context:** write it like a real person describing their job, so the reviewer can picture the workflow.

8. The 60-second check before you hit submit

1. Read your prompt out loud. Every “it/the/that” should point to something you actually named.
2. Find the two tools: where does it **read**, where does it **act**?
3. Find the dates and the timezone.
4. Find the finish line — the “done when...”
5. Make sure you didn’t write step 1, 2, 3.

If all five are there, submit.